



Deep Learning Project

COURSE CODE	CSE485
COURSE NAME	Deep Learning
SEMESTER	Fall 2024
DATE OF SUBMISSION	12/25/2024

Submitted By:

Kareem Mohamed Shoaib	20P5553
Adham Amr	20P5249
Mazen Tayseer	20P7460

Pre-Processing

```
numeric_columns = train_df.select_dtypes(include=['float64', 'int64']).columns
train_df[numeric_columns] = train_df[numeric_columns].fillna(train_df[numeric_columns].mean())

categorical_columns = train_df.select_dtypes(include=['object']).columns
for col in categorical_columns:
    train_df[col].fillna(train_df[col].mode()[0], inplace=True)
```

Handle Missing Values: Ensures that no NaN values remain in the dataset, which could otherwise cause errors during model training or data processing.

```
train_df.drop_duplicates(inplace=True)
```

Drop duplicates Removes duplicate rows from the train_df DataFrame. `inplace=True` ensures that the changes are applied directly to the train_df object without needing to reassign it.

```
label_encoder = LabelEncoder()
train_df['species'] = label_encoder.fit_transform(train_df['species'])
label_mapping = dict(zip(label_encoder.transform(label_encoder.classes_), label_encoder.classes_))
```

Encoding:

Encoding the categorical column species into numerical labels, which are easier for machine learning algorithms to process.

For example:

- 'Acer_Palmatum' → 0
- 'Quercus_Rubra' → 1
- 'Tilia_Tomentosa' → 2

Loading Images

```
def load_images(image_folder, image_names, target_size=(128, 128)):
    images = []
    for name in image_names:
        img_path = os.path.join(image_folder, name)
        img = cv2.imread(img_path)
        if img is not None:
            img = cv2.resize(img, target_size)
            images.append(img)
    return np.array(images)

image_folder = 'images'
train_images = load_images(image_folder, train_df['id'].astype(str) + '.jpg')
train_images = train_images / 255.0
```

Parameters:

- image_folder: The directory containing the images.
- image_names: A list of image file names to load (e.g., ['1.jpg', '2.jpg', ...]).
- target_size: The desired dimensions of the images after resizing (default is (128, 128)).

Process:

- Step 1: Initialize an empty list images to store loaded and processed images.
- Step 2: Iterate over each image file name in image_names:
 - Construct File Path: Combine image_folder and name using os.path.join to create the full file path for each image.
 - Load Image: Use cv2.imread() to load the image from the file path.
 - If the image is successfully loaded (not None):
 - Resize Image: Resize the image to the specified target_size using cv2.resize().
 - Append to List: Add the resized image to the images list.
- Step 3: Convert the images list into a NumPy array using np.array(). This ensures compatibility with TensorFlow and machine learning frameworks.

Train-Test-Split

```
train_df.drop('id', axis=1, inplace=True)

X_train, X_test, y_train, y_test = train_test_split(
    train_images, train_df['species'].values, test_size=0.2, stratify=train_df['species'].values, random_state=42
)

num_classes = len(np.unique(y_train))
```

□ **train_df.drop('id', axis=1, inplace=True):**

- Removes the id column from the train_df DataFrame because it is no longer required for training or testing.
- axis=1: Specifies that the column (not rows) should be dropped.
- inplace=True: Modifies the DataFrame directly without creating a copy.

□ **train_test_split Function:**

- Splits the dataset into training and testing sets.

Arguments:

- train_images: Array of preprocessed images (features).
- train_df['species'].values: The target labels (species) corresponding to each image.
- test_size=0.2: Specifies that 20% of the dataset will be used for testing, and the remaining 80% for training.
- stratify=train_df['species'].values:
 - Ensures that the proportion of each class (species) in the training and testing sets matches that of the original dataset.
 - This is important when the dataset has imbalanced classes.

Display Images

Display Sample Images

```
num_samples = 5
plt.figure(figsize=(15, 5))
for i in range(num_samples):
    plt.subplot(1, num_samples, i + 1)
    plt.imshow(X_train[i])
    plt.title(f"Label: {label_mapping[y_train[i]]}")
    plt.axis('off')
plt.show()
```

16]

✓ 0.3s

Python

Label: Quercus_Palustris



Label: Quercus_Rubra



Label: Zelkova_Serrata



Label: Cotinus_Coggygria



Label: Ginkgo_Biloba



Grid Search

```
param_grid = {
    'num_layers': [2],
    'dropout_rate': [0.5, 0.3],
    'optimizer_name': ['adam', 'sgd', 'rmsprop'],
    'weight_decay': [0.01, 0.001],
    'initial_lr': [0.001, 0.0001],
    'lr_scheduler': [
        None,
        'step_decay',
        'ReduceLROnPlateau'
    ],
}

param_combinations = list(product(
    param_grid['num_layers'],
    param_grid['dropout_rate'],
    param_grid['optimizer_name'],
    param_grid['weight_decay'],
    param_grid['initial_lr'],
    param_grid['lr_scheduler']
))

best_accuracy = 0
best_params = None
best_model_path = "leaf_classification_cnn_model.keras"
results = []
```

Training Function

```

    for params in param_combinations:
        num_layers, dropout_rate, optimizer_name, weight_decay, initial_lr, lr_scheduler = params
        print(f"\nTraining with params: {params}")

        model = Sequential()
        model.add(Conv2D(32, (3, 3), activation='relu', input_shape=(128, 128, 3),
                        kernel_regularizer=l2(weight_decay)))
        model.add(MaxPooling2D(pool_size=(2, 2)))

        for _ in range(num_layers - 1):
            model.add(Conv2D(64, (3, 3), activation='relu', kernel_regularizer=l2(weight_decay)))
            model.add(MaxPooling2D(pool_size=(2, 2)))

        model.add(Flatten())
        model.add(Dense(128, activation='relu', kernel_regularizer=l2(weight_decay)))
        model.add(Dropout(dropout_rate))
        model.add(Dense(y_train.max() + 1, activation='softmax', kernel_regularizer=l2(weight_decay)))

        if optimizer_name == 'adam':
            optimizer = Adam(learning_rate=initial_lr)
        elif optimizer_name == 'sgd':
            optimizer = SGD(learning_rate=initial_lr, momentum=0.9)
        elif optimizer_name == 'rmsprop':
            optimizer = RMSprop(learning_rate=initial_lr, decay=1e-6, rho=0.9)
        else:
            raise ValueError("Invalid optimizer. Choose 'adam', 'sgd', or 'rmsprop'.")

```

```

        callbacks.append(ReduceLROnPlateau(monitor='val_loss', factor=0.5, pat
    history = model.fit(
        X_train, y_train,
        validation_data=(X_test, y_test),
        epochs=20,
        batch_size=32,
        callbacks=callbacks,
        verbose=1
    )

```



```
history = model.fit(
    X_train, y_train,
    validation_data=(X_test, y_test),
    epochs=20,
    batch_size=32,
    callbacks=callbacks,
    verbose=1
)

_, test_accuracy = model.evaluate(X_test, y_test, verbose=0)
print(f"Test Accuracy: {test_accuracy * 100:.2f}%")

results.append((params, test_accuracy))

if test_accuracy > best_accuracy:
    best_accuracy = test_accuracy
    best_params = params
    model.save(best_model_path)
    print(f"New best model saved to {best_model_path}")

print("\nBest Parameters:")
print(best_params)
print(f"Best Accuracy: {best_accuracy * 100:.2f}%")
print(f"Best model saved to: {best_model_path}")

print("\nAll Results:")
for params, accuracy in results:
    print(f"Params: {params}, Accuracy: {accuracy * 100:.2f}%")
```


Results

Standardized data

```
Training with params: (2, 0.5, 'rmsprop', 0.01, 0.001, 'ReduceLROnPlateau')
Epoch 1/10
25/25 ————— 13s 397ms/step - accuracy: 0.0120 - loss: nan - val_accuracy: 0.0101 - val_loss: nan - learning_rate: 0.0010
Epoch 2/10
25/25 ————— 10s 378ms/step - accuracy: 0.0033 - loss: nan - val_accuracy: 0.0101 - val_loss: nan - learning_rate: 0.0010
Epoch 3/10
25/25 ————— 10s 392ms/step - accuracy: 0.0115 - loss: nan - val_accuracy: 0.0101 - val_loss: nan - learning_rate: 0.0010
Epoch 4/10
25/25 ————— 14s 551ms/step - accuracy: 0.0138 - loss: nan - val_accuracy: 0.0101 - val_loss: nan - learning_rate: 5.0000e-05
Epoch 5/10
25/25 ————— 12s 477ms/step - accuracy: 0.0069 - loss: nan - val_accuracy: 0.0101 - val_loss: nan - learning_rate: 5.0000e-05
Epoch 6/10
25/25 ————— 11s 458ms/step - accuracy: 0.0050 - loss: nan - val_accuracy: 0.0101 - val_loss: nan - learning_rate: 5.0000e-05
Epoch 7/10
25/25 ————— 12s 467ms/step - accuracy: 0.0099 - loss: nan - val_accuracy: 0.0101 - val_loss: nan - learning_rate: 2.5000e-05
Epoch 8/10
25/25 ————— 12s 477ms/step - accuracy: 0.0090 - loss: nan - val_accuracy: 0.0101 - val_loss: nan - learning_rate: 2.5000e-05
Epoch 9/10
25/25 ————— 12s 470ms/step - accuracy: 0.0074 - loss: nan - val_accuracy: 0.0101 - val_loss: nan - learning_rate: 2.5000e-05
Epoch 10/10
25/25 ————— 12s 492ms/step - accuracy: 0.0111 - loss: nan - val_accuracy: 0.0101 - val_loss: nan - learning_rate: 1.2500e-05
Test Accuracy: 1.01%
New best model saved to leaf_classification_cnn_model.keras
```

Accuracy with standardized data by computing the mean and standard deviation for each feature dimension using the training set only, then subtracting the mean and dividing by the stdev for each feature and each sample. Produced very bad accuracies and result , so we disposed it

Best Results after grid-search

```
Best Parameters:
(2, 0.5, 'rmsprop', 0.01, 0.001, 'step_decay')
Best Accuracy: 61.11%
Best model saved to: leaf_classification_cnn_model.keras
```

Evaluation

```
model_path = 'leaf_classification_cnn_model.keras'

model = tf.keras.models.load_model(model_path)
print(f"\nModel loaded from {model_path}")

model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])

print("\nEvaluating on Training Set:")
train_loss, train_accuracy = model.evaluate(X_train, y_train, verbose=0)
print(f"Training Accuracy: {train_accuracy * 100:.2f}%")

print("\nEvaluating on Test Set:")
test_loss, test_accuracy = model.evaluate(X_test, y_test, verbose=0)
print(f"Test Accuracy: {test_accuracy * 100:.2f}%")

y_pred = model.predict(X_test)
y_pred_classes = np.argmax(y_pred, axis=1)

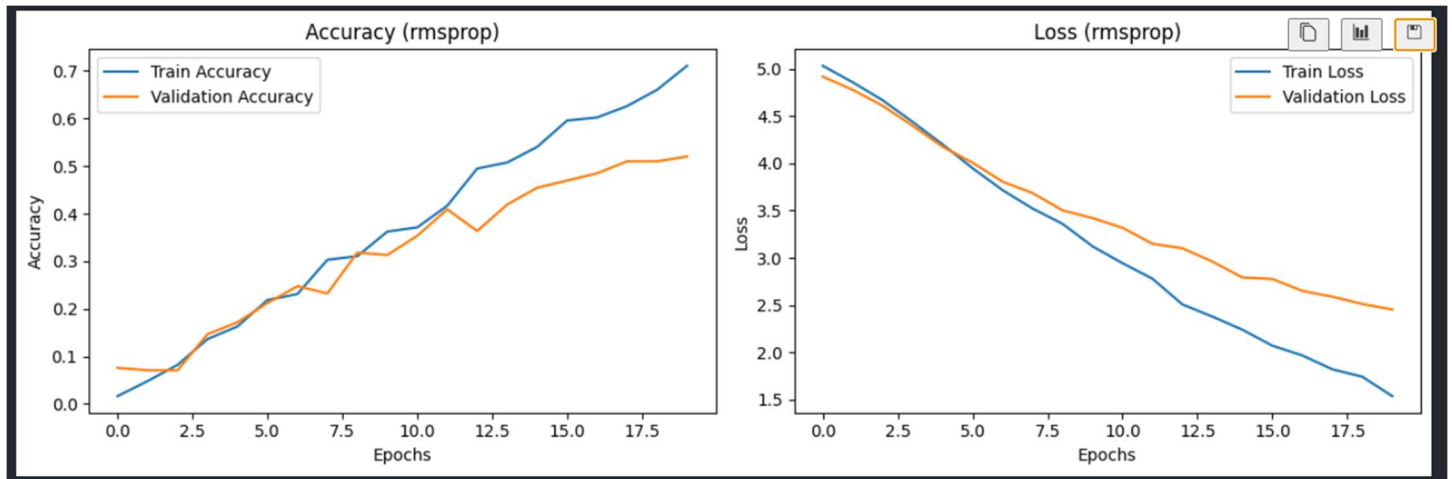
print("\nClassification Report (Test Set):")
print(classification_report(y_test, y_pred_classes))
```

```
Model loaded from leaf_classification_cnn_model.keras

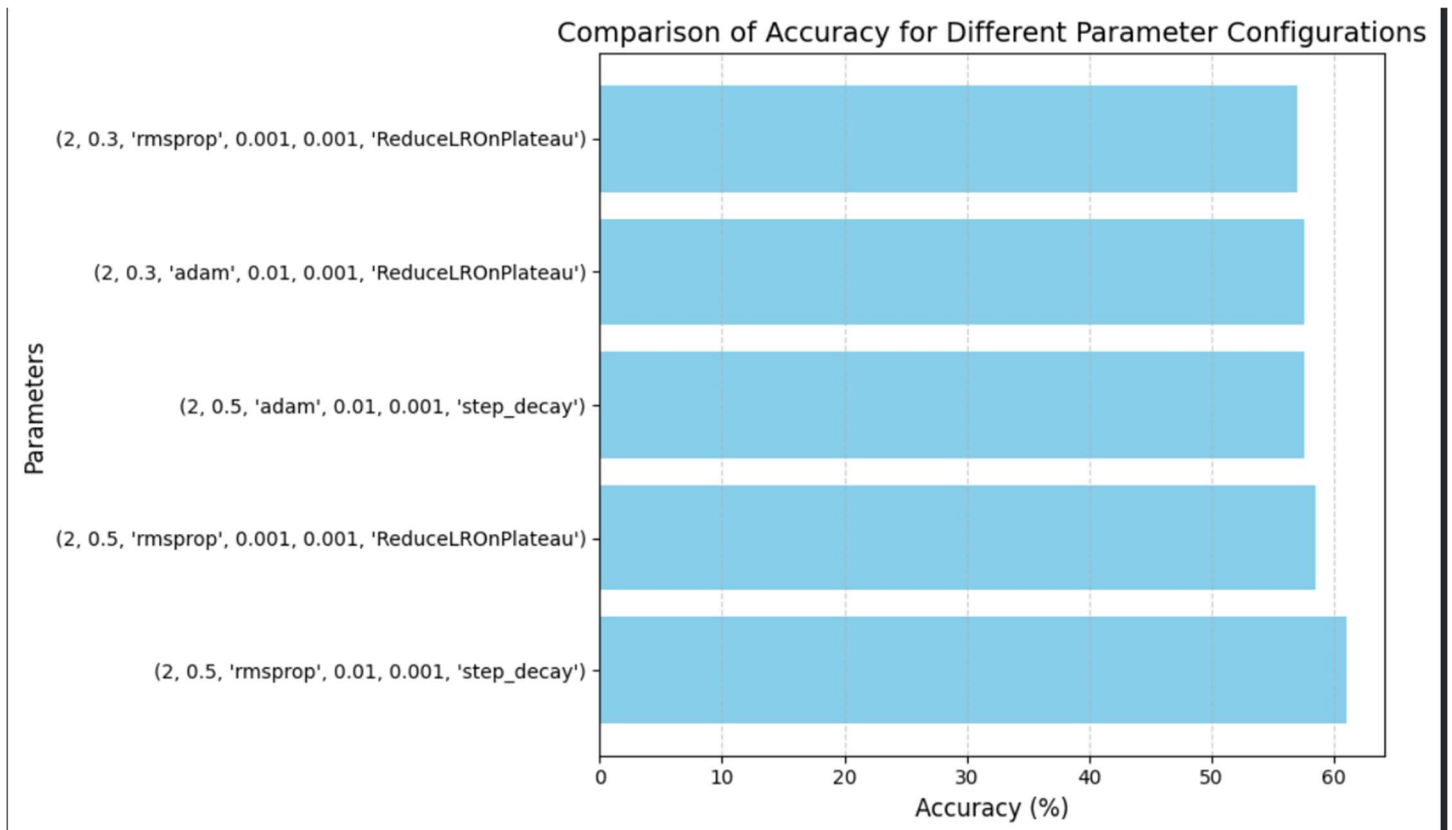
Evaluating on Training Set:
Training Accuracy: 98.23%

Evaluating on Test Set:
Test Accuracy: 61.11%
7/7 ————— 0s 41ms/step
```

Training history plot



Comparing with 5 best results



Classification of images

```
model = tf.keras.models.load_model(model_path)
print(f"\nModel loaded from {model_path}")

def preprocess_image(img_path, target_size=(128, 128)):
    img = cv2.imread(img_path)
    if img is not None:
        img = cv2.resize(img, target_size)
        img = img / 255.0
        return img
    return None

images = []
valid_image_names = []

test_images_folder = 'test_images'
test_images = ['1.jpg', '2.jpg', '3.jpg', '40.jpg']
for name in test_images:
    img_path = os.path.join(test_images_folder, name)
    img = preprocess_image(img_path)
    if img is not None:
        images.append(img)
        valid_image_names.append(name)
    else:
        print(f"Warning: Could not load image {name}. Skipping.")

if not images:
    print("No valid images to classify.")
else:
    images = np.array(images)
    predictions = model.predict(images)
    predicted_classes = np.argmax(predictions, axis=1)

    for i, name in enumerate(valid_image_names):
        print(f"Image: {name} -> Predicted Label: {label_mapping[predicted_classes[i]]}")
```

Model loaded from leaf_classification_cnn_model.keras

1/1 ————— 0s 203ms/step

Image: 1.jpg -> Predicted Label: Acer_Opalus

Image: 2.jpg -> Predicted Label: Pterocarya_Stenoptera

Image: 3.jpg -> Predicted Label: Quercus_Hartwissiana

Image: 40.jpg -> Predicted Label: Quercus_Phillyraeoides